

# Spring XD

## Introduction

Data Ingestion and Spring XD

# Introduction

- What is Data Ingestion, why do we need it?
- What is Spring XD?
- When to use Spring XD
- Lab

# What is Data Ingestion?

- Data ingestion is the process of:
  - reading, cleansing, transforming, and potentially enriching data from one or more sources;
  - and storing the resulting processed data into a consolidated repository



# Why Do We Need Data Ingestion?

- Modern applications generate massive data amounts.
- Modern data repositories are optimized for the task at hand
  - Not optimized for data analysis
- Data stored in diverse repositories difficult to analyze.
- Ingesting data into a consolidated repository permits practical data analysis.



# Data Ingestion Tools

- Various tools exist to simplify the data ingestion process
- Batch based:
  - Apache Sqoop
  - Spring XD
  - Spark
- Stream based:
  - Apache Flume
  - Spring XD
  - Spark



# Apache Spark

*Apache Spark is a fast and general engine for large-scale data processing.*

- <http://spark.apache.org/>

- **XD Streams**

- Ability to microbatch based on event count via Reactor and RxJava APIs.
- Ability to create data pipelines to process one event at a time.
- Flexibility to specify hosts to dictate the location of data computations.

- **XD Jobs**

- Provides REST-API and lifecycle management for Spark jobs.
- Extensible to integrate with other batch systems.

# Apache Flume

- *A service for collecting, aggregating, and moving large amounts of log data.*
  - <http://flume.apache.org>

- Spring XD uses an interactive shell and DSL for stream creation, while Flume uses property (key/value pair) files.
- Administration and monitoring via the admin UI.
- Granular controls to manifest batch job and step execution to create complex data driven workflows.
- Flexibility through a deployment manifest to declar- atively configure data partitioning strategy to route
- data to a specific consumer instance in the cluster.

# Apache Sqoop

*Tool for transferring bulk data between Apache Hadoop and structured datastores (such as relational databases.)*

- <http://sqoop.apache.org/>

- Orchestrate Pig, Hive, HBase, MapReduce or other batch systems.
- Extend and customize batch workflow infrastructure.
- High level configuration DSL to create, deploy and destroy batch workflows.
- Operationalize custom data pipelines through REST.
- Unified functional programming support to build reactive-style data pipelines.



# Apache Oozie

*Oozie is a workflow scheduler engine to manage Hadoop workloads such as MapReduce or Pig jobs.*

- <http://sqoop.apache.org/>

- Building upon Spring Batch, a JSR standardization (JSR-352) of batch workload data processing, Spring XD inherits workflow scheduling and execution functionalities.
- Provides out of the box batch jobs that support plain files, JDBC, HDFS, FTP, MongoDB, Spark and Sqoop.
- Ability to scale jobs without having to bring down the runtime.
- Provides bi-directionality between real-time streaming and batch workflows to accommodate complex data processing use cases.
- Ability to create, and launch jobs from the admin UI. • Ability to view historical snapshots of job executions from the admin UI.

# Spring XD

*"Spring XD is a unified, distributed, and extensible service for data ingestion, real time analytics, batch processing, and data export."*

- <http://docs.spring.io/spring-xd/docs/1.1.0.RELEASE/reference/html/>

# Spring XD

## Problems

**Batch and Streaming are often handled by multiple platforms**

**Fragmented Big Data Ecosystem**

**Not all data is Hadoop bound**

## Spring XD Benefits

### **Unified Approach - Stream Processing and Batch Jobs**

- Hadoop Batch workflow orchestration
- Analytics
- Machine Learning Scoring
- Eye on the big picture

### **Runtime provides critical Non-functional requirements**

- Scalable, Distributed, Fault-Tolerant
- Portable on prem DIY cluster, YARN, EC2, (WIP for PCF)
- Easy to use, extend and integrate other technologies

### **Proven**

- Built on proven EAI and Batch spring projects (7years)

# Spring XD



- Project home: <http://projects.spring.io/spring-xd/>
- "XD": "Extreme Data"
- *Lambda* Architecture
- Support for both stream processing and batch jobs
- Scalable, distributed, and fault-tolerant
- Deploy to
  - On premise cluster
  - YARN
  - EC2

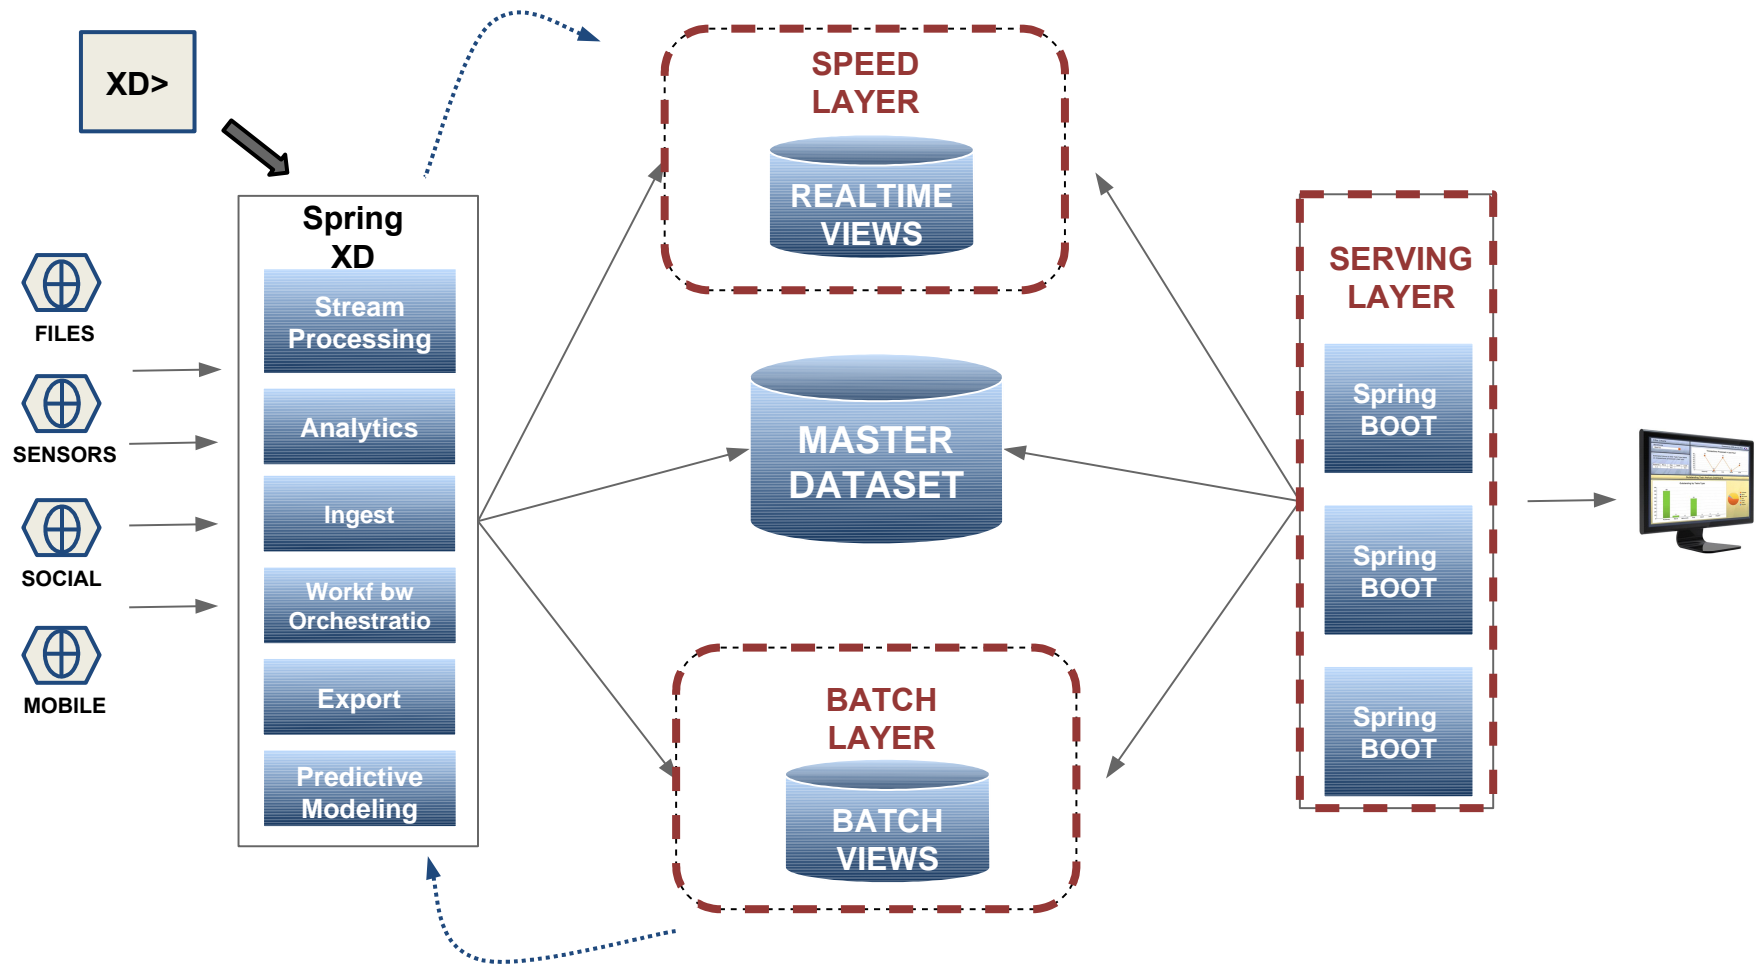
# Lambda Architecture

*"Data processing architecture designed to handle massive quantities of data by taking advantage of both batch and stream processing methods."*


- 3 layers
  - Batch: Pre-computes and stores results from all available data
  - Speed: Processes data streams in real time
  - Serving: Stores output from batch and speed layers, amalgamates results from both into views

- [http://en.wikipedia.org/wiki/Lambda\\_architecture](http://en.wikipedia.org/wiki/Lambda_architecture)

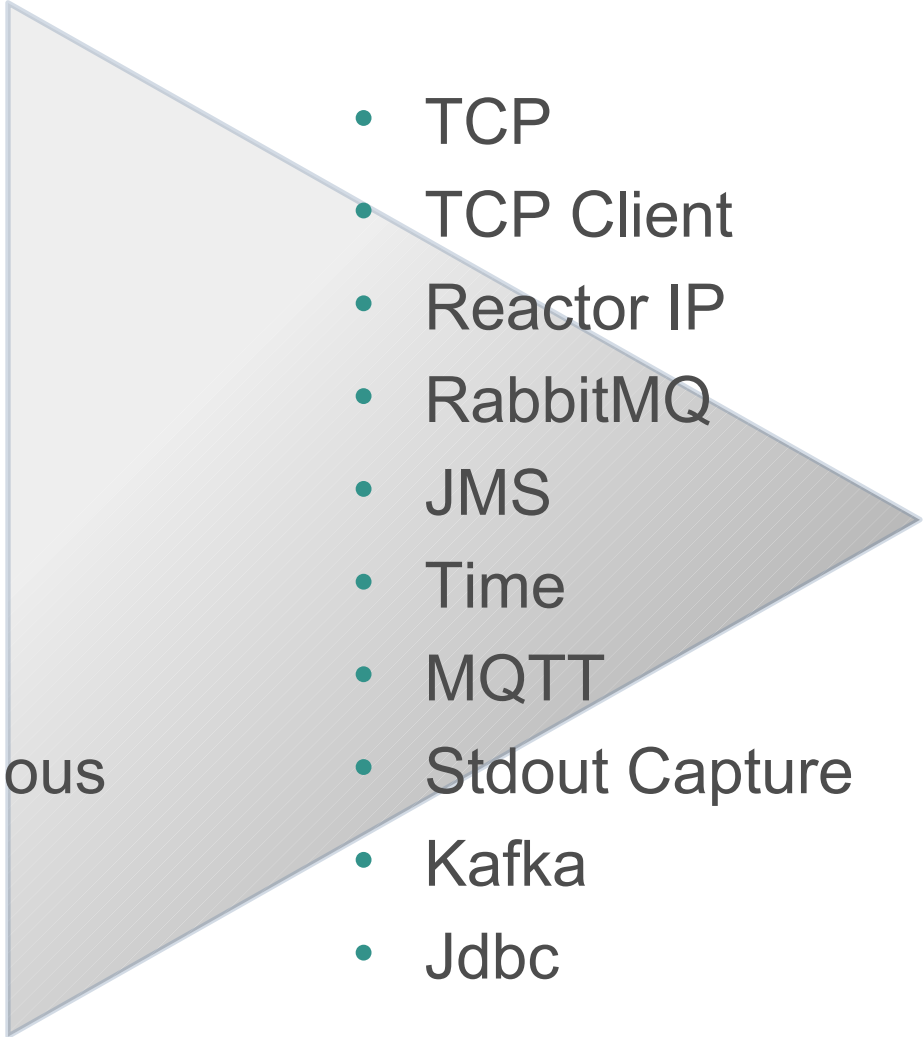
# Lambda Architecture



# Streams

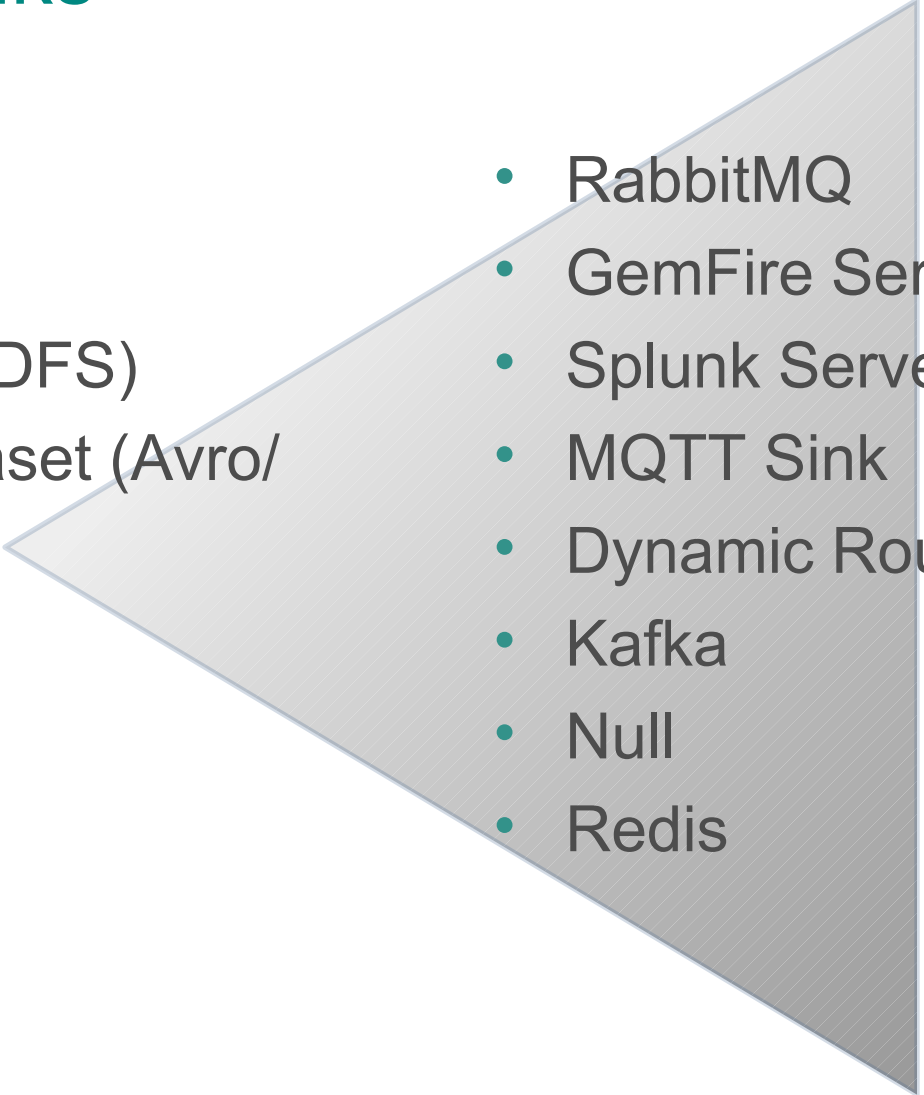


# Input Sources

- 
- HTTP
  - Tail
  - File
  - Mail
  - Twitter Search
  - Twitter Stream
  - Gemfire Source
  - Gemfire Continuous Query
  - Syslog
  - TCP
  - TCP Client
  - Reactor IP
  - RabbitMQ
  - JMS
  - Time
  - MQTT
  - Stdout Capture
  - Kafka
  - Jdbc



# Output Sinks

- 
- Log
  - File Sink
  - Hadoop (HDFS)
  - HDFS Dataset (Avro/Parquet)
  - JDBC
  - TCP Sink
  - Mongo
  - Mail
  - RabbitMQ
  - GemFire Server
  - Splunk Server
  - MQTT Sink
  - Dynamic Router
  - Kafka
  - Null
  - Redis

# Stream Processing

- Filter
- Transform
- Script
- Splitter
- Aggregator
- HTTP Client
- JSON to Tuple
- Object to JSON

# Batch Jobs

- Support for batch processing includes:
  - Import CSV files to HDFS
  - Import CSV files to JDBC
  - Import JDBC to HDFS
  - Export HDFS to JDBC
  - Export HDFS to MongoDB
  - Spark Job

# When To Use Spring XD

- High throughput distributed data ingestion from various input sources into big data stores
- Real-time analytics at ingestion time
- Workflow management via batch jobs
- High throughput data export

# Use Cases

Telco service provider continuously **analyzing customer usage patterns** to up-sell products and services. Perform **analytics to predict churn** and take corrective actions to avoid customer acquisition costs.

Retail service provider delivering personalized mobile content including targeted advertisements and promotional offers based on geo-location.

Agriculture service provider creating batch workflows to govern pre, current, post and delivery phases of harvesting lifecycle.

e-Commerce service provider delivering customized, personalized and lucrative offers to online shoppers in real-time.

# Getting Ready for the Labs

- Let's setup \$LAB\_HOME
  - \$LAB\_HOME == spring-xd-one-day directory
- Git
  - git clone <https://github.com/cppwfs/spring-xd-one-day.git>
- Thumb Drive
  - Copy the spring-xd-one-day directory to your local drive

# LAB #1 Installation

- 1) Go to the instructions section of your \$LABS\_DIR.
- 2) Open either the lab/index.html or lab.pdf files
- 3) Follow the instructions on Chapter 1
  - Git
    - `git clone https://github.com/cppwfs/spring-xd-one-day.git`