

Spring XD

Advanced Streams

Advanced Streams

- **Monitoring**
- Hadoop (HDFS)
- Transform/Script processors
- TCP/UDP sources
- Composite Modules
- Lab

Monitoring Deployed Streams

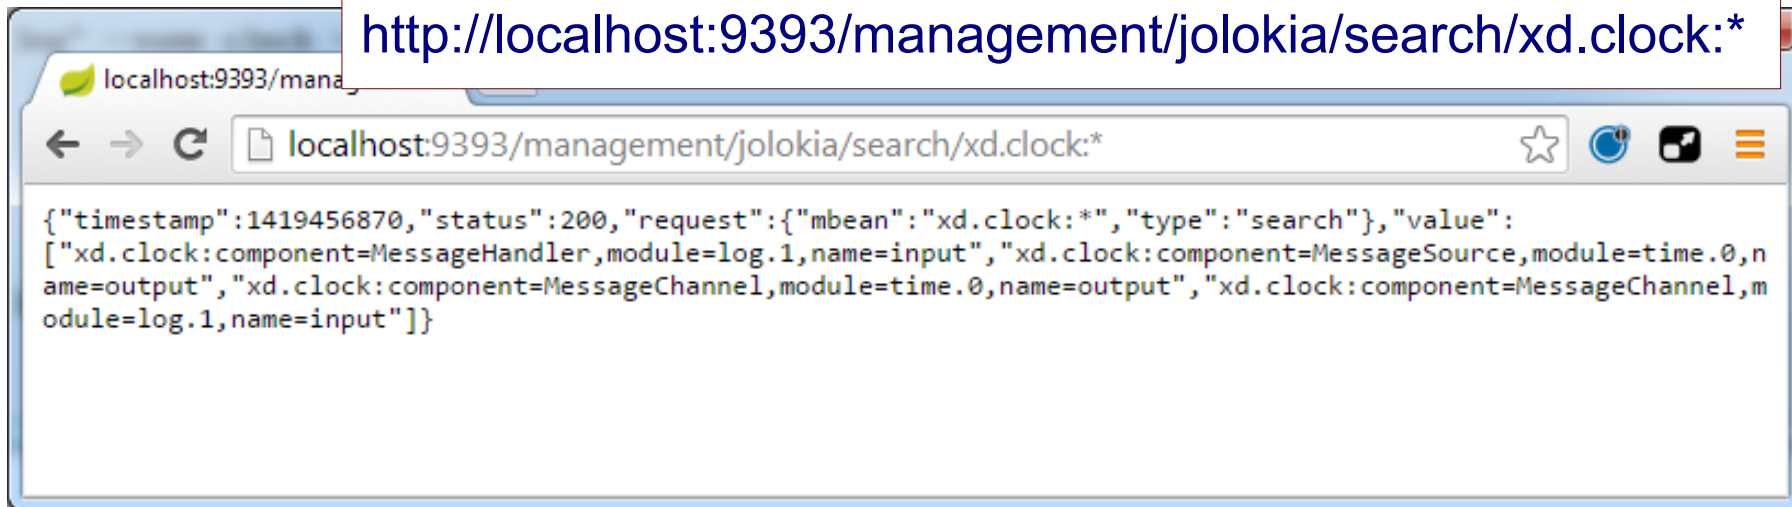
- Example: Fetch all modules for deployed stream “clock”

```
xd:>stream create --definition "time | log" --name clock  
Created new stream 'clock'  
xd:>stream deploy --name clock  
Deployed stream 'clock'
```

- Creates Mbeans with domain name xd.clock
- Creates objects time.0 and log.1

Monitoring Deployed Streams

http://localhost:9393/management/jolokia/search/xd.clock:*



```
{ "timestamp": 1419456870, "status": 200, "request": {
  "mbean": "xd.clock:*", "type": "search", "value": [
    "xd.clock:component=MessageHandler,module=log.1,name=input",
    "xd.clock:component=MessageSource,module=time.0,name=output",
    "xd.clock:component=MessageChannel,module=time.0,name=output",
    "xd.clock:component=MessageChannel,module=log.1,name=input"
  ]
}
```

Stream Attributes and Operations

- Source modules have `MessageSourceMetrics` attributes and operations available
- Processor, sink modules have `MessageHandlerMetrics` attributes and operations available
- Channels have `MessageChannelMetrics` attributes and operations available


Hadoop Support

- Major Hadoop distributions are supported
- Needed to be specified during startup of containers and shell using `--hadoopDistro` argument
- See next slide for supported values or execute `"xd-shell --hadoopDistro"`

```
$ xd-container --hadoopDistro phd21
$ xd-shell --hadoopDistro phd21
$ xd-singlenode --hadoopDistro phd21
```

Supported Hadoop Distributions

- Supported values for `--hadoopDistro` arg:
 - `hadoop25` - Apache Hadoop 2.5.2
 - `hadoop26` - Apache Hadoop 2.6.0 (default)
 - `phd1` - Pivotal HD 1.1
 - `phd21` - Pivotal HD 2.1
 - `phd21` - Pivotal HD 2.1
 - `cdh5` - Cloudera CDH 5.3
 - `hdp22` - Hortonworks Data Platform 2.2



List varies with
Spring XD versions

Hadoop Setup

- Hadoop client not required
 - No core-site.xml
 - Nameserver config in xd/config/servers.yml (default localhost:8020)

```
# Hadoop properties
spring:
  hadoop:
    fsUri: hdfs://192.168.218.140:8020
```

- Additional config in xd/config/hadoop.properties

HDFS Sink

- Writes raw data to one or more HDFS files
 - Simple example:

```
stream create
  --name myTimeHdfs
  --definition "time | hdfs"
  --deploy
```

- Creates the file /xd/myTimeHdfs/myTimeHdfs-0.txt in HDFS.
- But, file might not be closed for long time
 - hence, not readable from other processes

HDFS Open Files

- Spring XD adds a prefix/suffix to open files
 - Default suffix: `.tmp`
 - Can be changed with `--inUseSuffix` and `--inUsePrefix` options

```
stream create
  --name myTimeHdfs
  --definition
    "time | hdfs --inUseSuffix=' ' --inUsePrefix='tmp- '"
  --deploy
```

HDFS Rollover

- Default rollover is 1GB
 - can be changed with the rollover option:

```
stream create
  --name myTimeHdfs
  --definition "time | hdfs --rollover=256K"
  --deploy
```

- Value is treated as bytes by default.
 - Other examples: 64K, 128M, 512G, 1T
- Forces early closing of file, results in file names like:
 - myTimeHdfs-0.txt, myTimeHdfs-1.txt, myTimeHdfs-2.txt

HDFS Idle Time

- Rollover only good for events with continuous high rate
- Not so good for low frequency burst events
 - E.g. from HTTP Source. No consistent state in closed files on HDFS.
- Use `idleTimeout` instead (milliseconds)

```
stream create
  --name myTimeHdfs
  --definition "http | hdfs --idleTimeout=3000"
  --deploy
```

- Can be combined with rollover

HDFS Naming Strategies

- Default file location is
 - `/xd/<stream name>/<stream name>-<rolling part>.txt`
- Can be changed using
 - `--directory`, `--fileName` and `--fileExtension`

```
stream create
  --name myTimeHdfs
  --definition
    "time | hdfs ---directory=/my --fileName=data"
  --deploy
```

- Creates files in `/my/data-0.txt`, `/my/data-1.txt`, ...

XD Shell HDFS commands

- Spring XD shell supports some useful HDFS commands
 - No need to install Hadoop client!
- Needs to be configured first:

```
xd:> hadoop config fs --namenode hdfs://localhost:8020
```

- Supports basic HDFS operations, like:

```
xd:> hadoop fs ls /xd
xd:> hadoop fs cat /xd/myTimeHdfs/myTimeHdfs-0.txt
xd:> hadoop fs rm /xd/myTimeHdfs/myTimeHdfs-0.txt
xd:> hadoop fs rm /xd --recursive
```

Advanced Streams

- Hadoop (HDFS)
- **Transform/Script processors**
- TCP/UDP sources
- Composite Modules
- Lab

Transform Processor

- Generic way to convert messages using a Spring Integration transformer
- Supports Spring Expression Language (SpEL) expressions and Groovy scripts

```
xd:> stream create --name myTransform --definition
"http | transform --expression='payload.toUpperCase()' '
| log" --deploy
Created and deployed new stream 'myTransform'
xd:>http post --data "Transform me!"
> POST (text/plain;Charset=UTF-8) http://localhost:9000
Transform me!
> 200 OK
```

SpEL

```
13:30:34,535 1.1.0.RELEASE INFO pool-21-thread-4
sink.myTransform - TRANSFORM ME!
```


Groovy Script Locations

- Groovy scripts need to be placed in `${xd.home}/modules/processor/scripts`
- Or specify alternate location using `--script` option and `file:` prefix
 - Recommend using a shared NFS location

```
xd:> stream create --name myTransform --definition
"http
| transform --script=file:/myscripts/myUpperCase.groovy
| log" --deploy
```

Transform Processor

- Example with Groovy script
- Script *must* return a value

```
return payload.toUpperCase()
```

```
xd:> stream create --name myTransform --definition  
"http  
| transform --script=myUpperCase.groovy  
| log" --deploy
```



```
$ curl -d "Test" http://localhost:9000
```

```
19:07:47,105 1.0.2.RELEASE INFO pool-31-thread-4  
sink.myTransform - TEST
```

Script Processor

- Generic way to process messages using a Spring Integration Service Activator
- Script does *not* have to return a value

`return null`

```
xd:> stream create --name myTransform --definition
"http
| script --script=myUpperCase.groovy
| log" -deploy
Created and deployed new stream 'myTransform'
```

```
xd:>http post --data "Drop me!"
> POST (text/plain;Charset=UTF-8) http://localhost:9000 Drop
me!
> 200 OK
```

Difference Between Filter/Transform/Script

- All of them can execute Groovy scripts
- Filter: Groovy scripts decides if message gets dropped
 - Return false to drop message
- Transform: Groovy script needs to return something, can only convert message
 - Throws an exception if no value is returned
- Script: Groovy script can convert message or drop
 - Return null to drop message
- Conclusion: Script is the most generic one

Advanced Streams

- Hadoop (HDFS)
- Transform/Script processors
- **TCP/UDP sources**
- Composite Modules
- Lab

TCP Source

- Acts as a raw TCP server
- Several decoders available to frame a message out of the TCP stream
 - CRLF (default), LF, NULL, STXEX, ..
- Produces byte[], can be changed with **--outputType=text/plain**

```
xd:> stream create --name tcptest --definition  
    "tcp --port=3000 --outputType=text/plain | log"  
    --deploy
```

TCP Source Decoders (Framing)

```
xd:> stream create --name tcptest --definition  
      "tcp --port=3000 --outputType=text/plain | log"  
      --deploy
```

Default
CRLF Decoder

```
$ echo -en 'foobar\r\n' | netcat localhost 3000
```

Send <CRLF> terminated
string to localhost:3000

```
xd:> stream create --name tcptest --definition  
      "tcp --port=3000 --outputType=text/plain --decoder=NULL  
      | log" --deploy
```

null Decoder

```
$ echo -en 'foobar\x00' | netcat localhost 3000
```

Send null terminated
string to localhost:3000

Using Spring XD to Capture Stdout

- No stdin source provided
- Use TCP source instead:

```
xd:> stream create --name tcpforstdout --definition  
      "tcp --port=3000 --outputType=text/plain --decoder=LF  
      | log" --deploy
```

```
$ cat mylog.txt | netcat localhost 3000
```


Advanced Streams

- Hadoop (HDFS)
- Transform/Script processors
- TCP/UDP sources
- **Composite Modules**
- Lab

Composite Modules

- Possible to combine 2 or more modules into a composite module
- Promotes re-use
- Improves performance
 - Composite module is deployed as a single unit
 - Local in-memory channel will be used instead of messaging middleware
 - Fewer network hops

Composite Modules

- Given the following streams:

```
stream1 = http | filter
           --expression=payload.contains('foo') | file
stream2 = file | filter
           --expression=payload.contains('foo') | file
```



A yellow box with the word "Identical" in black text is positioned to the right of the code snippets. Two lines extend from the box, one pointing to the filter expression in the first snippet and the other pointing to the filter expression in the second snippet, indicating that the two snippets are identical.

- Create a composite module from the common components

```
xd:> module compose foo
      --definition "filter
      --expression=payload.contains('foo') | file"
```

Composite Modules

- Composite modules will appear in the `module list` with a `(c)` beside them
- Create the streams using the composite component

```
xd:> stream create httpfoo --definition  
      "http | foo" --deploy
```

```
xd:> stream create filefoo --definition "file  
      --outputType=text/plain | foo" --deploy
```

Reactor IP Source

- Comparable to TCP Source, but based on Reactor project
 - Much higher message throughput
 - Also supports UDP (`--transport=udp`)
 - Produces Strings as default (can be changed with `--codec`)
- Only supports linefeed (default) and length framing
 - Called decoders in TCP source

```
xd:> stream create --name tcptest --definition  
      "reactor-ip --port=3000 | log" --deploy
```

```
$ echo -en 'foobar\n' | netcat localhost 3000
```

Composite Modules

- A composite module that contains:
 - a source and one or more processors is considered a source module
 - only processors is considered a processor module
 - a sink and one or more processors is considered a sink module

- In the example below we are creating the foo sink

```
xd:> module compose foo --definition "filter  
--expression=payload.contains('foo') | file"
```

- To delete the module you would have to refer to it as:

```
xd:> module delete -name sink:foo
```

Lab

Experiment with Transformers and Composed Modules